

Day 1

Introduction to probabilistic programming languages (PPLs)

Andrés Masegosa

*Department of Mathematics
University of Almería
Spain*

This material jointly made with Thomas D. Nielsen (AAU).

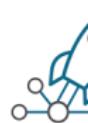
Day 1: Introduction to probabilistic programming languages

- Why do we need PPLs?
- Probabilistic programming in Pyro
- Hand-on exercises:
 - Probability Distributions in Pyro.
 - Probabilistic Models in Pyro.
 - Ice-cream Shop Model.

Day 2: Probabilistic Models with Deep Neural Networks

- Uncertainty in Machine Learning
- Variational Inference
- Supervised/Unsupervised Learning
- Hand-on exercises
 - Bayesian Neural Networks
 - Variational Auto-encoders

Why PPLs?



PyMC3



Why do we need PPLs?



PyMC3



Why do we need PPLs?

- Reason 1: “Democratization” of the development of AI systems.

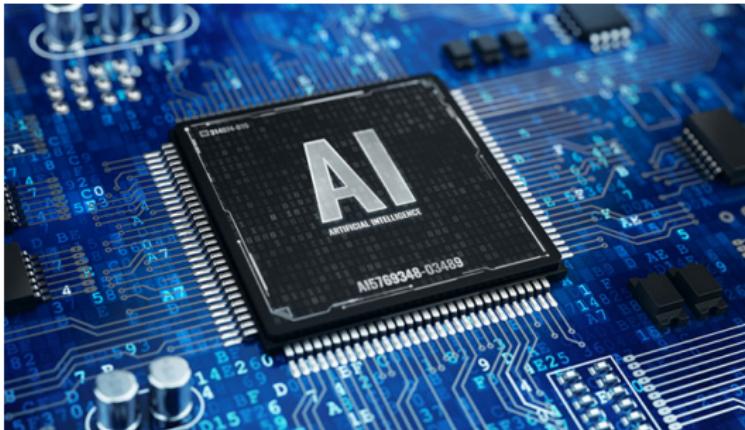


Why do we need PPLs?

- **Reason 1:** “Democratization” of the development of AI systems.
- **Reason 2:** Make AI systems safer.

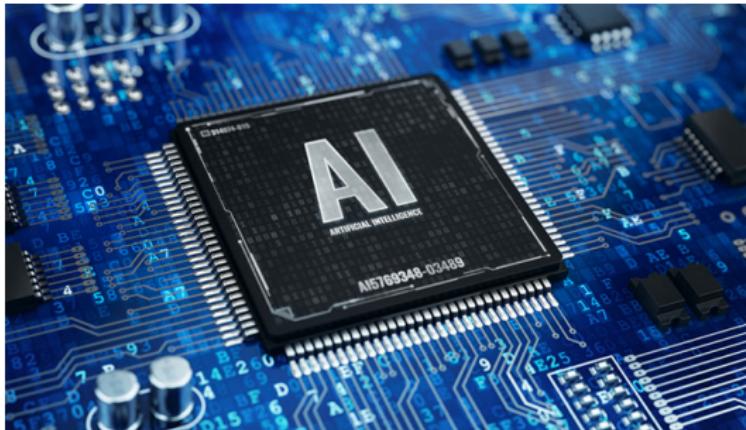
Reason 1: “Democratization” of the development of AI systems.

Why PPLs?



The development of machine learning systems requires enormous efforts.

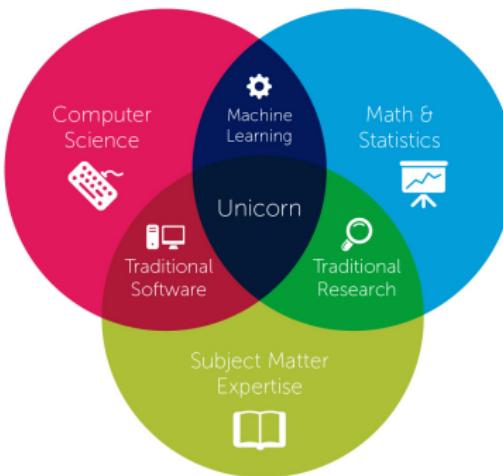
Why PPLs?



The development of machine learning systems requires enormous efforts.

DARPA's Fund Call for PPLs in Artificial Intelligence.

Data Science

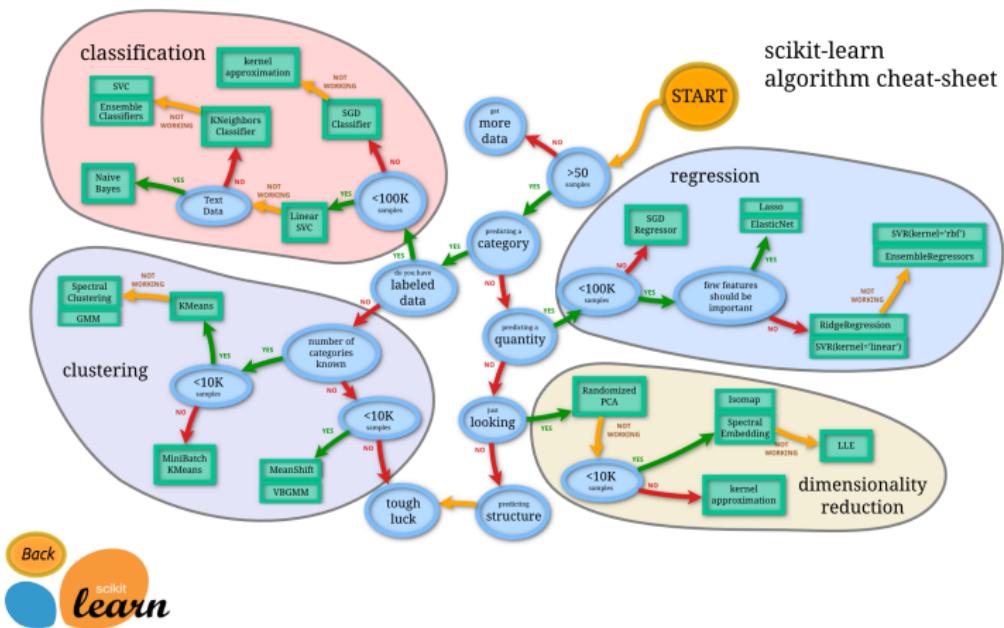


Copyright © 2014 by Steven Geringer Raleigh, NC.
Permission is granted to use, distribute, or modify this image,
provided that this copyright notice remains intact.

The development of **machine learning systems** requires enormous efforts.

- It requires of highly qualified experts.

Why PPLs?

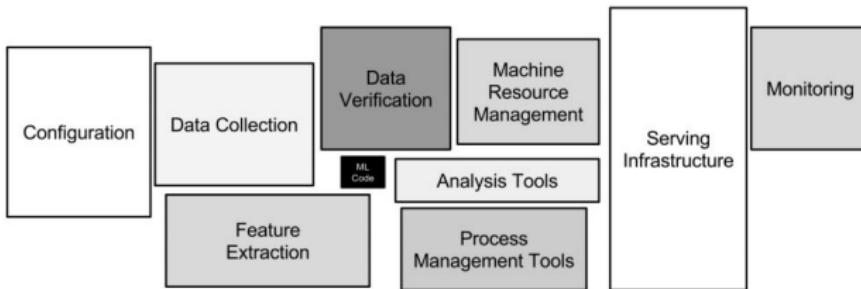


The development of machine learning applications requires enormous effort.

- It is necessary to have highly qualified experts.
- **It is difficult to find the ML model most suitable for an application.**

Hidden Technical Debt in Machine Learning Systems

D. Sculley, Gary Holt, Daniel Golovin, Eugene Davydov, Todd Phillips
{dsculley, gholt, dg, edavydov, toddphillips}@google.com
Google, Inc.



The development of machine learning applications requires enormous effort.

- It is necessary to have highly qualified experts.
- It is difficult to find the ML model most suitable for an application.
- **Programming a ML model is a complex task where many problems are intermingled.**

Wanted: Artificial intelligence experts

In artificial intelligence, job openings are rising faster than job seekers.



Consequences:

- Shortage of AI experts (and high salaries).

Wanted: Artificial intelligence experts

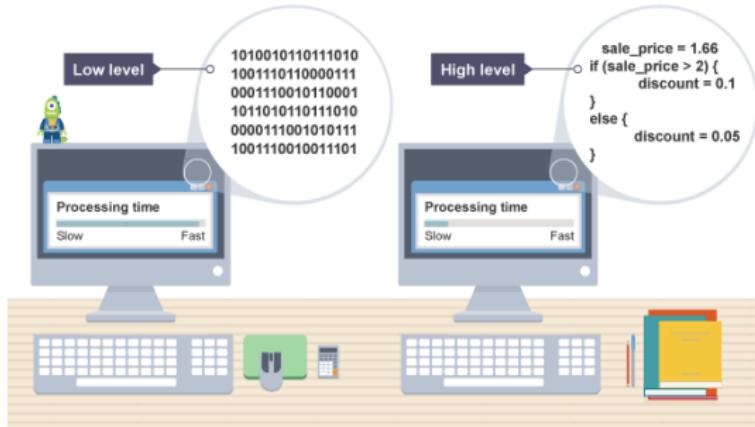
In artificial intelligence, job openings are rising faster than job seekers.



Consequences:

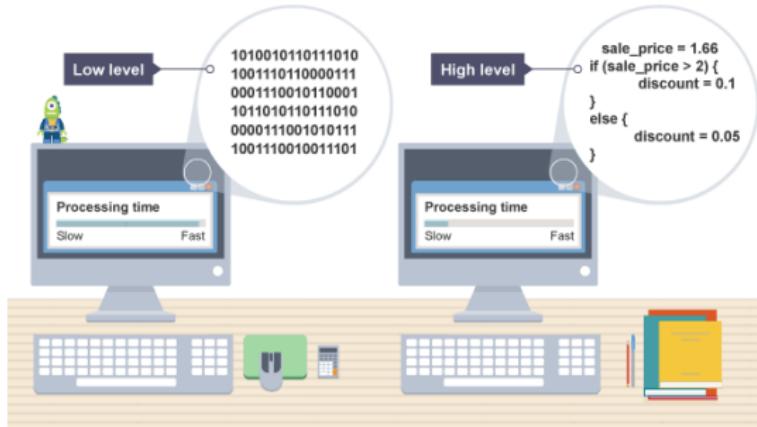
- Shortage of AI experts (and high salaries).
- Only big corporations have the resources for developing ML systems.

Why PPLs?



Similar situation than 50 years ago:

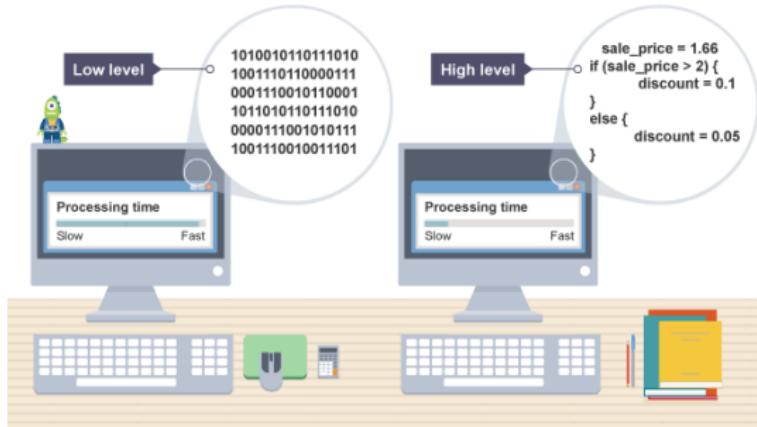
Why PPLs?



Similar situation than 50 years ago:

- People used to program in low-level programming languages.

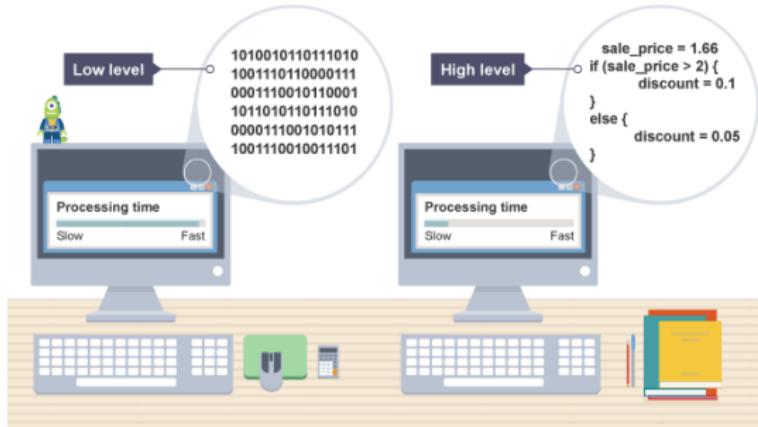
Why PPLs?



Similar situation than 50 years ago:

- People used to program in low-level programming languages.
- Programming was complex and demand high-expertise.

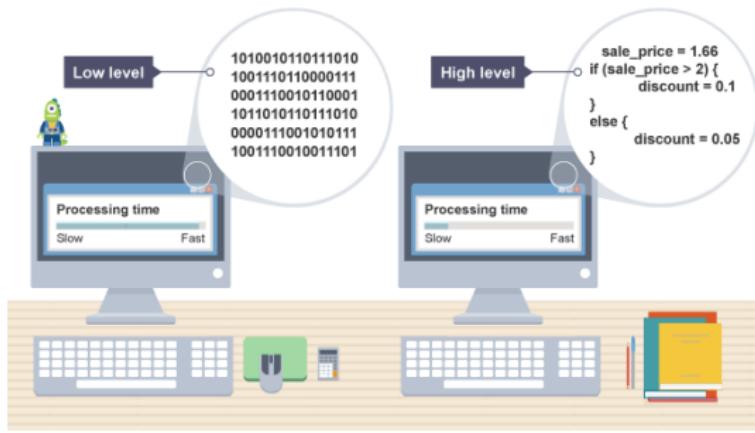
Why PPLs?



Similar situation than 50 years ago:

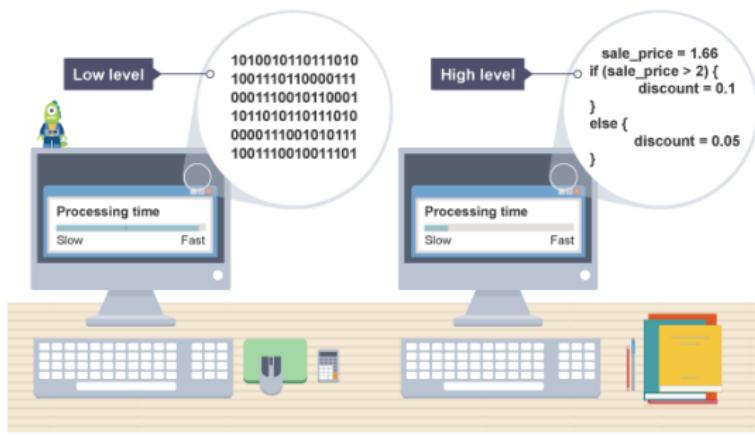
- People used to program in low-level programming languages.
- Programming was complex and demand high-expertise.
- Focus on application and low-level hardware details.

Why PPLs?



High-level programming languages brought many advantages:

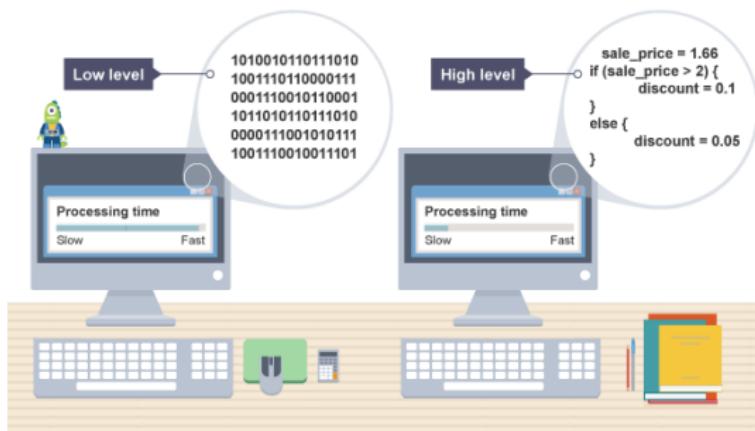
Why PPLs?



High-level programming languages brought many advantages:

- Programmers focused on the applications.

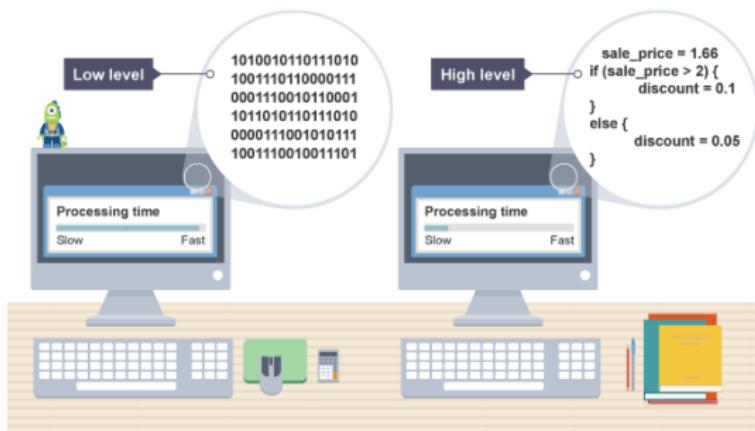
Why PPLs?



High-level programming languages brought many advantages:

- Programmers focused on the applications.
- Hardware Experts focused on compilers.

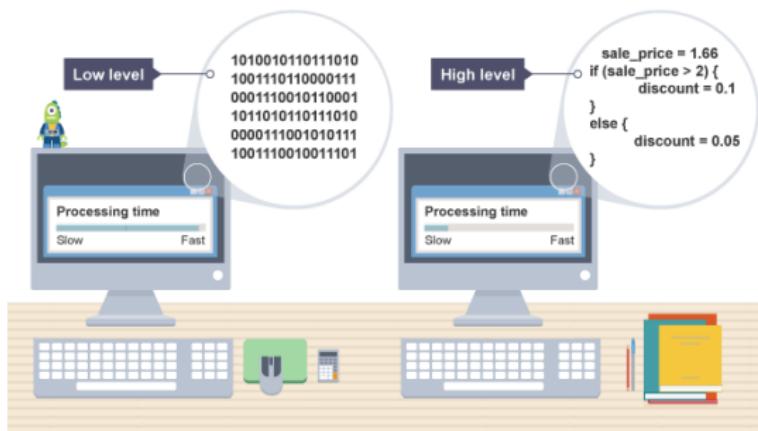
Why PPLs?



High-level programming languages brought many advantages:

- Programmers focused on the applications.
- Hardware Experts focused on compilers.
- High gains in productivity.

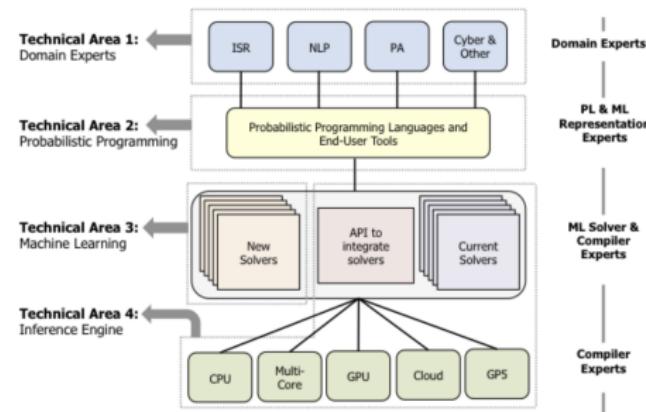
Why PPLs?



High-level programming languages brought many advantages:

- Programmers focused on the applications.
- Hardware Experts focused on compilers.
- High gains in productivity.
- “Democratization” of the software development.

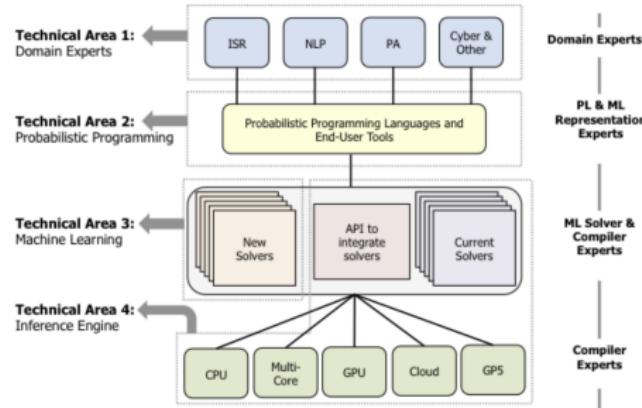
Why PPLs?



PPLs as high-level programming languages for **machine learning systems**:

- Stacked architecture

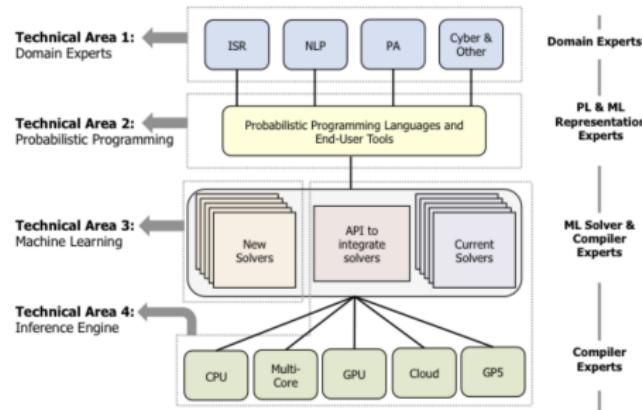
Why PPLs?



PPLs as high-level programming languages for **machine learning systems**:

- Stacked architecture
- Different Domain Experts will code their models using the same language.

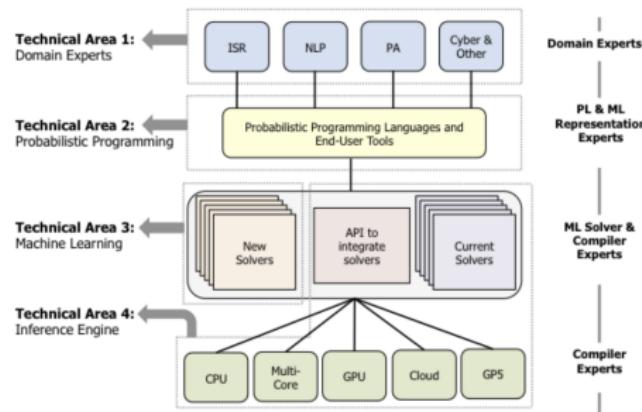
Why PPLs?



PPLs as high-level programming languages for **machine learning systems**:

- Stacked architecture
- Different Domain Experts will code their models using the same language.
- ML experts will focus on the development of new ML solvers.

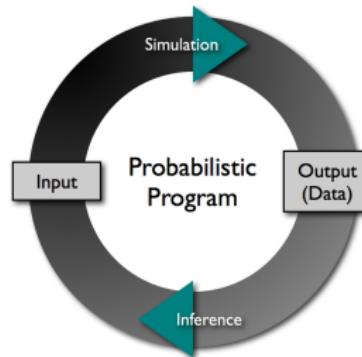
Why PPLs?



PPLs as high-level programming languages for **machine learning systems**:

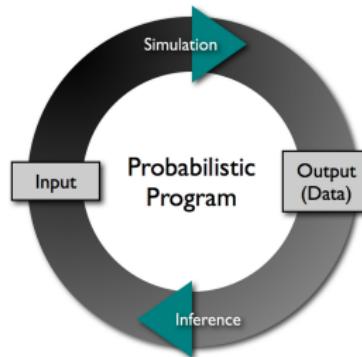
- Stacked architecture
- Different Domain Experts will code their models using the same language.
- ML experts will focus on the development of new ML solvers.
- Compile experts will focus on running these ML solvers on specialized hardware.

Why PPLs?



Benefits of PPLs:

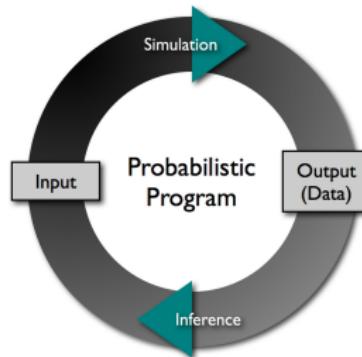
- Simplify machine learning model code.



Benefits of PPLs:

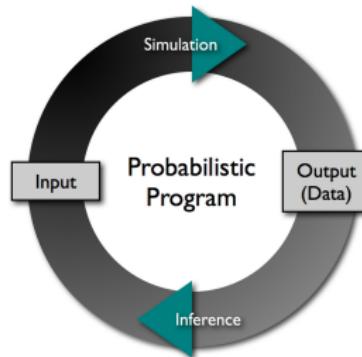
- Simplify machine learning model code.
- Reduce development time and cost to encourage experimentation.

Why PPLs?



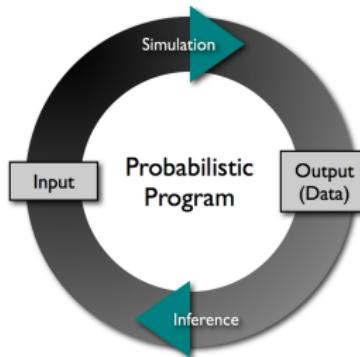
Benefits of PPLs:

- Simplify machine learning model code.
- Reduce development time and cost to encourage experimentation.
- Facilitate the construction of more sophisticated models.



Benefits of PPLs:

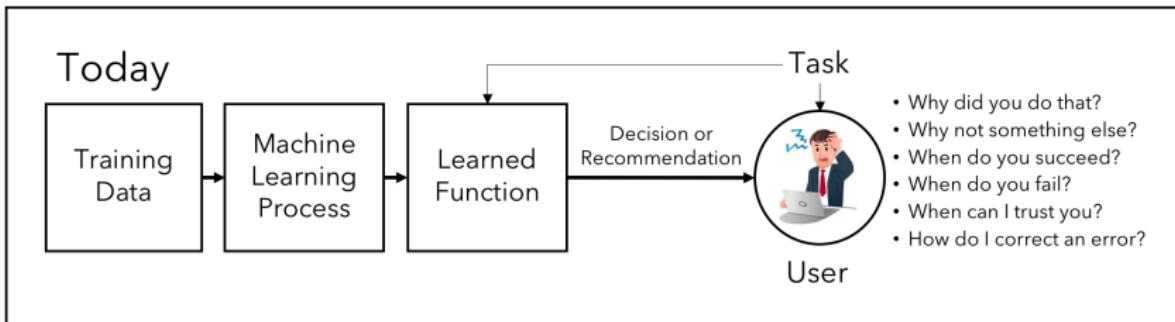
- Simplify machine learning model code.
- Reduce development time and cost to encourage experimentation.
- Facilitate the construction of more sophisticated models.
- Reduce the necessary level of expertise.



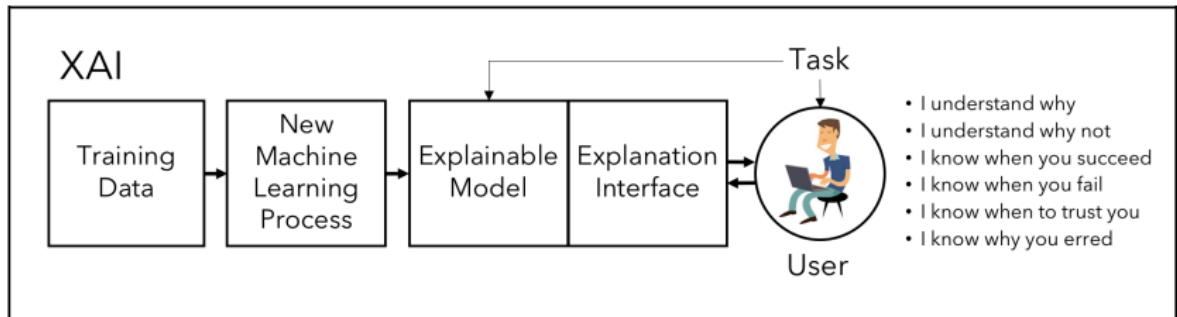
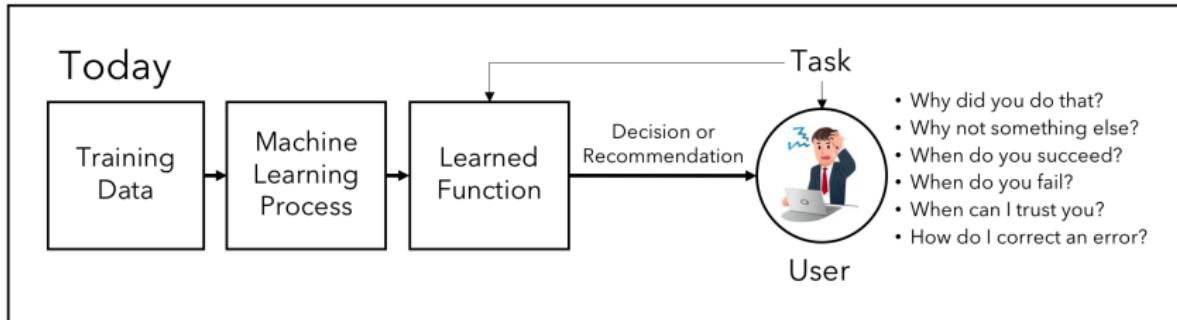
Benefits of PPLs:

- Simplify machine learning model code.
- Reduce development time and cost to encourage experimentation.
- Facilitate the construction of more sophisticated models.
- Reduce the necessary level of expertise.
- “Democratization” of the development of ML systems.

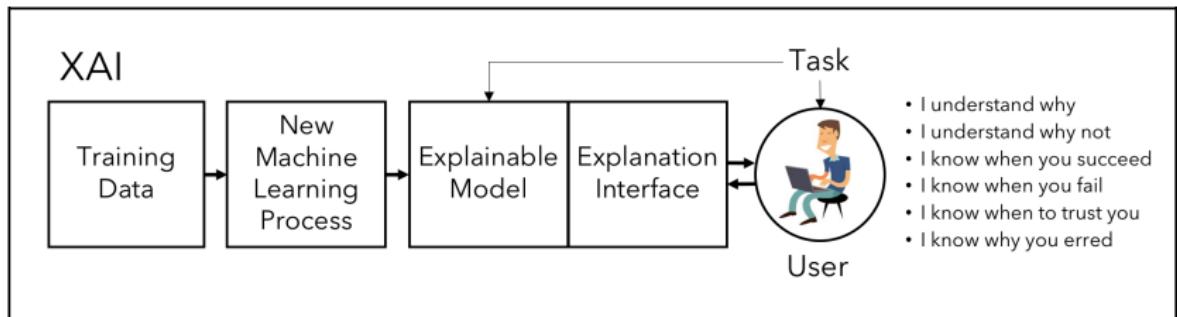
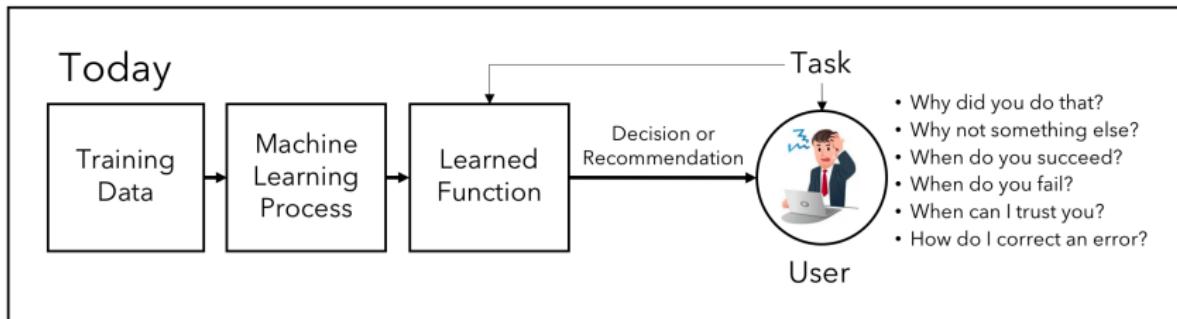
Reason 2: Make AI systems safer



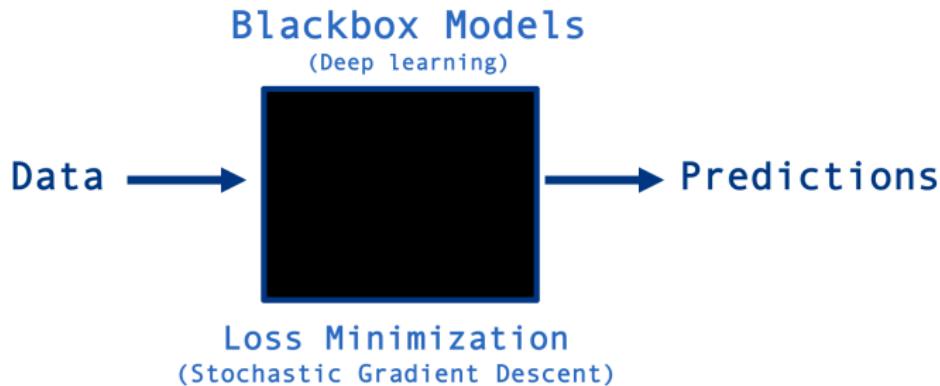
XAI: Explainable AI



XAI: Explainable AI

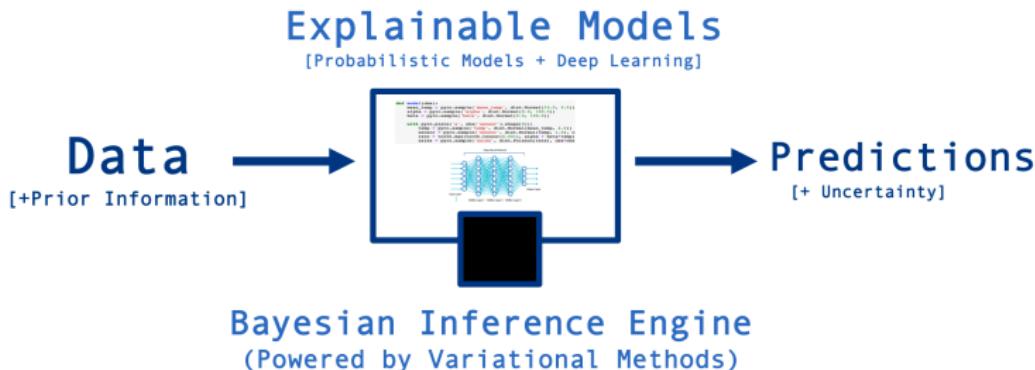


DARPA's Fund Call for XAI projects.



Deep Learning Based AI Systems :

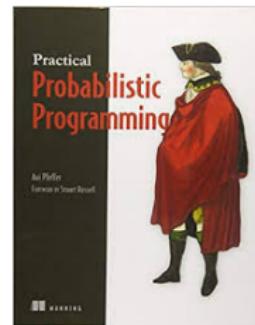
- BlackBox Models impossible to **interpret**.
- Impossible to know **how sure they are** in their predictions.
- Enormously **limit** the application of AI to many **real life problems**.



PPL based Systems :

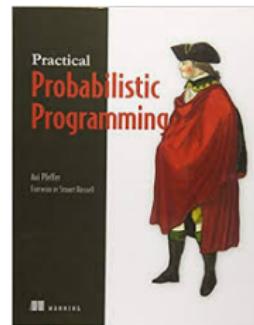
- PPLs provide **transparent** model description.
- PPLs provide **uncertainty estimation** both in models and predictions.
- PPLs aim to combine the **best of deep learning and probabilistic models**.

Brief Historical Review of PPLs



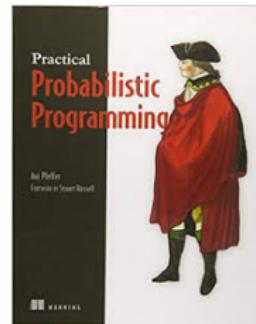
1st Generation of PPLs :

- Bugs, WinBugs, Jags, Figaro, etc.



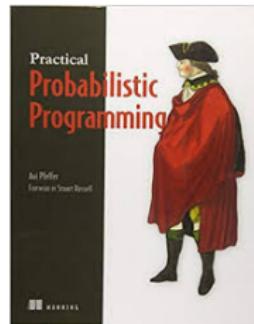
1st Generation of PPLs :

- Bugs, WinBugs, Jags, Figaro, etc.
- **Turing-complete** probabilistic programming languages. (i.e. they can represent any computable probability distribution).



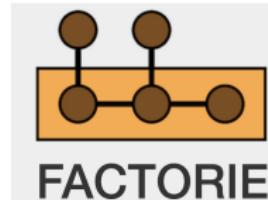
1st Generation of PPLs :

- Bugs, WinBugs, Jags, Figaro, etc.
- **Turing-complete** probabilistic programming languages. (i.e. they can represent any computable probability distribution).
- Inference engine based on **Monte Carlo** methods.



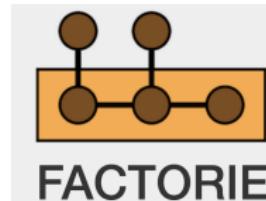
1st Generation of PPLs :

- Bugs, WinBugs, Jags, Figaro, etc.
- **Turing-complete** probabilistic programming languages. (i.e. they can represent any computable probability distribution).
- Inference engine based on **Monte Carlo** methods.
- They did **not scale** to large data samples/high-dimensional models.



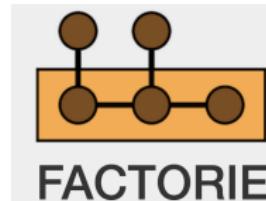
2nd Generation of PPLs :

- Infer.net, Factorie, Amidst, etc.



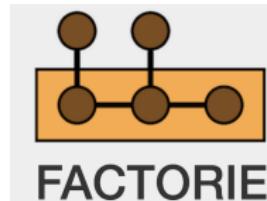
2nd Generation of PPLs :

- Infer.net, Factorie, Amidst, etc.
- Inference engine based on **message passage algorithms** and/or variational inference methods.



2nd Generation of PPLs :

- Infer.net, Factorie, Amidst, etc.
- Inference engine based on **message passage algorithms** and/or variational inference methods.
- They did **scale** to large data samples/high-dimensional models.



2nd Generation of PPLs :

- Infer.net, Factorie, Amidst, etc.
- Inference engine based on **message passage algorithms** and/or variational inference methods.
- They did **scale** to large data samples/high-dimensional models.
- **Restricted** probabilistic model family (i.e. factor graphs, conjugate exponential family, etc.)



PyMC3



3rd Generation of PPLs :

- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.



PyMC3



3rd Generation of PPLs :

- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.
- **Black Box Variational Inference** and Hamiltonian Monte-Carlo.



PyMC3



3rd Generation of PPLs :

- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.
- **Black Box Variational Inference** and Hamiltonian Monte-Carlo.
- They did **scale** to large data samples/high-dimensional models.



PyMC3



3rd Generation of PPLs :

- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.
- **Black Box Variational Inference** and Hamiltonian Monte-Carlo.
- They did **scale** to large data samples/high-dimensional models.
- **Turing-complete** probabilistic programming languages.



PyMC3



3rd Generation of PPLs :

- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.
- **Black Box Variational Inference** and Hamiltonian Monte-Carlo.
- They did **scale** to large data samples/high-dimensional models.
- **Turing-complete** probabilistic programming languages.
- Rely on **deep learning frameworks** (TensorFlow, Pytorch, Theano, etc).



3rd Generation of PPLs :

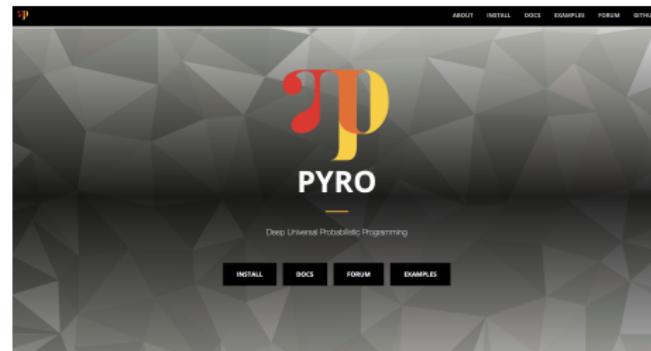
- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.
- **Black Box Variational Inference** and Hamiltonian Monte-Carlo.
- They did **scale** to large data samples/high-dimensional models.
- **Turing-complete** probabilistic programming languages.
- Rely on **deep learning frameworks** (TensorFlow, Pytorch, Theano, etc).
 - Specialized hardware like GPUs, TPUs, etc.

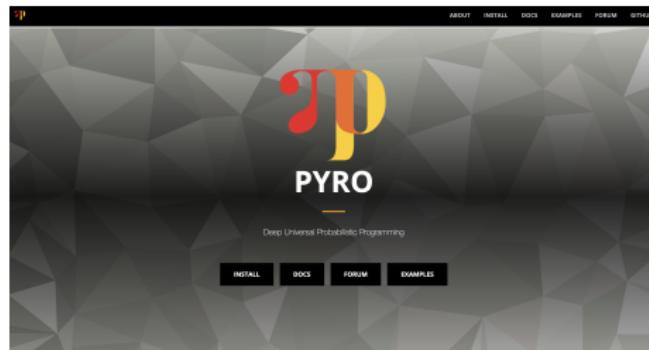


3rd Generation of PPLs :

- TensorFlow Probability, Pyro, PyMC3, InferPy, etc.
- **Black Box Variational Inference** and Hamiltonian Monte-Carlo.
- They did **scale** to large data samples/high-dimensional models.
- **Turing-complete** probabilistic programming languages.
- Rely on **deep learning frameworks** (TensorFlow, Pytorch, Theano, etc).
 - Specialized hardware like GPUs, TPUs, etc.
 - Automatic differentiation methods.

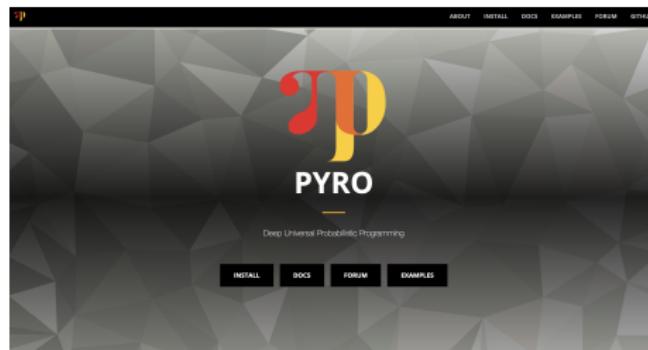
Pyro





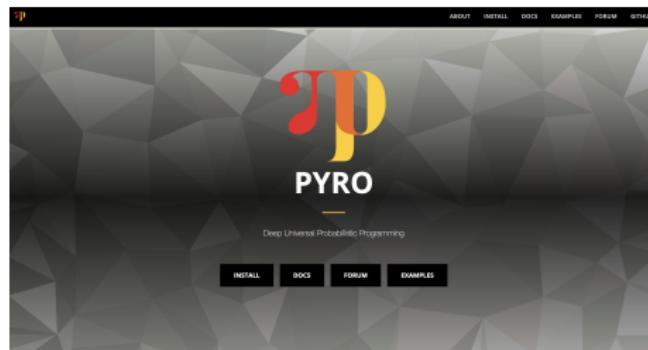
Pyro's main features (www.pyro.ai) :

- Developed by **UBER** (the car riding company).



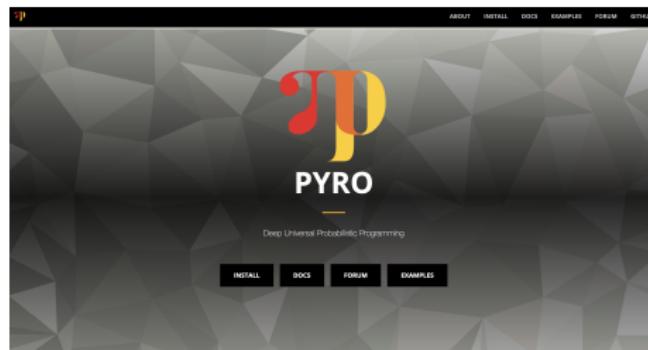
Pyro's main features (www.pyro.ai) :

- Developed by **UBER** (the car riding company).
- Focus on **probabilistic models with deep neural networks**.



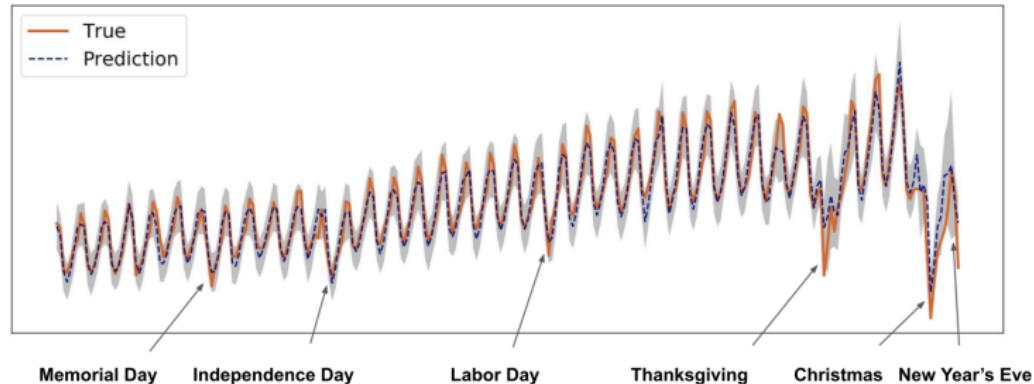
Pyro's main features (www.pyro.ai) :

- Developed by **UBER** (the car riding company).
- Focus on **probabilistic models with deep neural networks**.
- Rely on **Pytorch** (Deep Learning Framework).



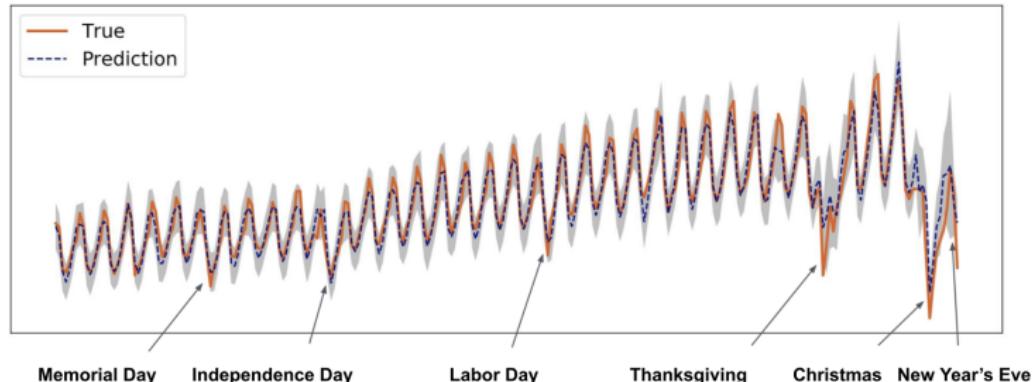
Pyro's main features (www.pyro.ai) :

- Developed by **UBER** (the car riding company).
- Focus on **probabilistic models with deep neural networks**.
- Rely on **Pytorch** (Deep Learning Framework).
- Enable **GPU** acceleration and distributed learning.



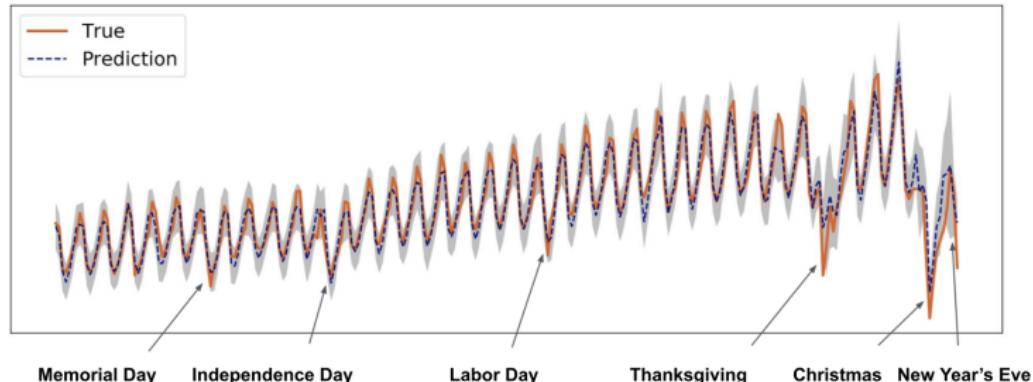
Demand Prediction with Pyro:

- Demand prediction is **critical** for user experience, resource allocation, etc.



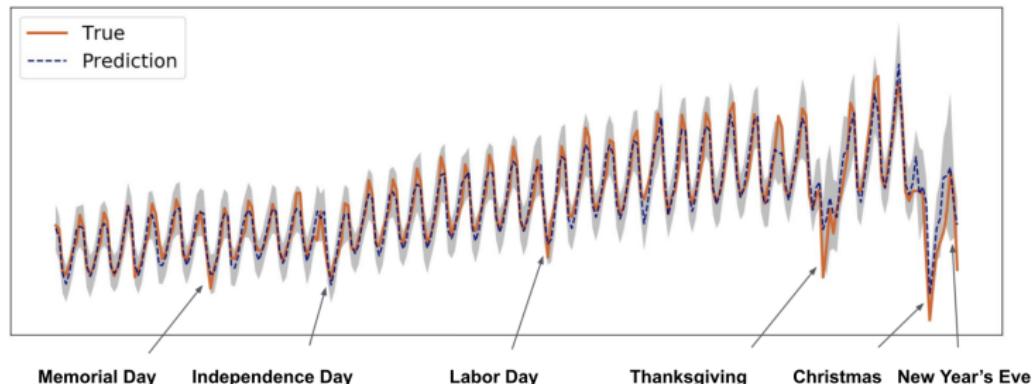
Demand Prediction with Pyro:

- Demand prediction is **critical** for user experience, resource allocation, etc.
- Prediction at **special events** is challenging: weather, population growth, etc.



Demand Prediction with Pyro:

- Demand prediction is **critical** for user experience, resource allocation, etc.
- Prediction at **special events** is challenging: weather, population growth, etc.
- LSTM powerful for time series modelling, but do not capture the **inherent variance**.



Demand Prediction with Pyro:

- Demand prediction is **critical** for user experience, resource allocation, etc.
- Prediction at **special events** is challenging: weather, population growth, etc.
- LSTM powerful for time series modelling, but do not capture the **inherent variance**.
- Bayesian LSTM provides **uncertainty estimation**.

Pyro's Distributions

```
In [3]: normal = dist.Normal(0,1)  
normal
```

```
Out[3]: Normal(loc: 0.0, scale: 1.0)
```

Pyro's distributions (<http://docs.pyro.ai/en/stable/distributions.html>) :

- Wide range of distributions: Normal, Beta, Cauchy, Dirichlet, Gumbel, Poisson, Pareto, etc.

```
In [15]: sample = normal.sample()
sample
Out[15]: tensor(0.4908)
```

Pyro's distributions (<http://docs.pyro.ai/en/stable/distributions.html>) :

- Wide range of distributions: Normal, Beta, Cauchy, Dirichlet, Gumbel, Poisson, Pareto, etc.
- Samples from the distributions are Pytorch's Tensor objects (i.e. multidimensional arrays).

```
In [17]: sample = normal.sample(sample_shape=[3,4,5])
sample.shape
```



```
Out[17]: torch.Size([3, 4, 5])
```

Pyro's distributions (<http://docs.pyro.ai/en/stable/distributions.html>) :

- Wide range of distributions: Normal, Beta, Cauchy, Dirichlet, Gumbel, Poisson, Pareto, etc.
- Samples from the distributions are Pytorch's Tensor objects (i.e. multidimensional arrays).

```
In [19]: torch.sum(normal.log_prob(sample))
```

```
Out[19]: tensor(-85.1003)
```

Pyro's distributions (<http://docs.pyro.ai/en/stable/distributions.html>) :

- Wide range of distributions: Normal, Beta, Cauchy, Dirichlet, Gumbel, Poisson, Pareto, etc.
- Samples from the distributions are Pytorch's Tensor objects (i.e. multidimensional arrays).
- Operations, like log-likelihood, are defined over tensors (with GPU acceleration powered by Pytorch).

```
In [9]: normal = dist.Normal(torch.tensor([1.,2.,3.]),1.)  
normal
```

```
Out[9]: Normal(loc: torch.Size([3]), scale: torch.Size([3]))
```

Pyro's distributions (<http://docs.pyro.ai/en/stable/distributions.html>) :

- Wide range of distributions: Normal, Beta, Cauchy, Dirichlet, Gumbel, Poisson, Pareto, etc.
- Samples from the distributions are Pytorch's Tensor objects (i.e. multidimensional arrays).
- Operations, like log-likelihood, are defined over tensors (with GPU acceleration powered by Pytorch).
- Multiple distributions can be embedded in single object (to define efficient vectorized operations).

```
▶ In [10]: normal.sample()
```

```
Out[10]: tensor([2.0592, 2.4035, 3.1918])
```

Pyro's distributions (<http://docs.pyro.ai/en/stable/distributions.html>) :

- Wide range of distributions: Normal, Beta, Cauchy, Dirichlet, Gumbel, Poisson, Pareto, etc.
- Samples from the distributions are Pytorch's Tensor objects (i.e. multidimensional arrays).
- Operations, like log-likelihood, are defined over tensors (with GPU acceleration powered by Pytorch).
- Multiple distributions can be embedded in single object (to define efficient vectorized operations).

```
In [11]: normal.log_prob(normal.sample())
```

```
Out[11]: tensor([-0.9402, -1.2113, -2.3214])
```

Pyro's distributions (<http://docs.pyro.ai/en/stable/distributions.html>) :

- Wide range of distributions: Normal, Beta, Cauchy, Dirichlet, Gumbel, Poisson, Pareto, etc.
- Samples from the distributions are Pytorch's Tensor objects (i.e. multidimensional arrays).
- Operations, like log-likelihood, are defined over tensors (with GPU acceleration powered by Pytorch).
- Multiple distributions can be embedded in single object (to define efficient vectorized operations).

Open the Notebook and Play around

- Test that you have installed the basic packages.
- Test that you can run the first lines of code.
- Play a bit with the code in Section 1 of the notebook.

`Day1/students_PPLs_Intro.ipynb`

Pyro's Models

```
In [12]: def model():
    temp = pyro.sample('temp', dist.Normal(15.0, 2.0))
    return temp

print(model())
print(model())

tensor(12.3926)
tensor(22.5272)
```

Pyro's models (http://pyro.ai/examples/intro_part_i.html) :

- A probabilistic model is defined as a **stochastic function**.

```
In [12]: def model():
    temp = pyro.sample('temp', dist.Normal(15.0, 2.0))
    return temp

print(model())
print(model())

tensor(12.3926)
tensor(22.5272)
```

Pyro's models (http://pyro.ai/examples/intro_part_i.html) :

- A probabilistic model is defined as a **stochastic function**.
- Each random variable is associated to a **primitive stochastic function** using the construct **pyro.sample(...)**.

```
In [21]: def model():
    temp = pyro.sample('temp', dist.Normal(15.0, 2.0))
    sensor = pyro.sample('sensor', dist.Normal(temp, 1.0))
    return (temp, sensor)

out1 = model()
out1

Out[21]: (tensor(15.8576), tensor(16.9907))
```

Pyro's models (http://pyro.ai/examples/intro_part_i.html) :

- A **stochastic function** can be defined as a composition of **primitive stochastic functions**.

```
In [21]: def model():
    temp = pyro.sample('temp', dist.Normal(15.0, 2.0))
    sensor = pyro.sample('sensor', dist.Normal(temp, 1.0))
    return (temp, sensor)

out1 = model()
out1

Out[21]: (tensor(15.8576), tensor(16.9907))
```

Pyro's models (http://pyro.ai/examples/intro_part_i.html) :

- A **stochastic function** can be defined as a composition of **primitive stochastic functions**.
- We define the **joint probability distribution**:

$$p(\text{sensor}, \text{temp}) = p(\text{sensor}|\text{temp})p(\text{temp})$$

Pyro's Inference

```
In [22]: #The observations
obs = {'sensor': torch.tensor(18.0)}

def model(obs):
    temp = pyro.sample('temp', dist.Normal(15.0, 2.0))
    sensor = pyro.sample('sensor', dist.Normal(temp, 1.0), obs=obs['sensor'])
```

Pyro's inference (http://pyro.ai/examples/intro_part_ii.html) :

- We can introduce observations (e.g. $\text{sensor} = 18.0$).

```
In [27]: #Run inference
svi(model,guide,obs, plot=True)

#Print results
print("P(Temperature|Sensor=18.0) = ")
print(dist.Normal(pyro.param("mean").item(), pyro.param("scale").item()))
print("")
```

Pyro's inference (http://pyro.ai/examples/intro_part_ii.html) :

- We can introduce observations (e.g. sensor = 18.0).
- We can query the **posterior probability distribution**:

$$p(\text{temp}|\text{sensor} = 18) = \frac{p(\text{sensor} = 18|\text{temp})p(\text{temp})}{\int p(\text{sensor} = 18|\text{temp})p(\text{temp})d\text{temp}}$$

```
In [27]: #Run inference
svi(model,guide,obs, plot=True)

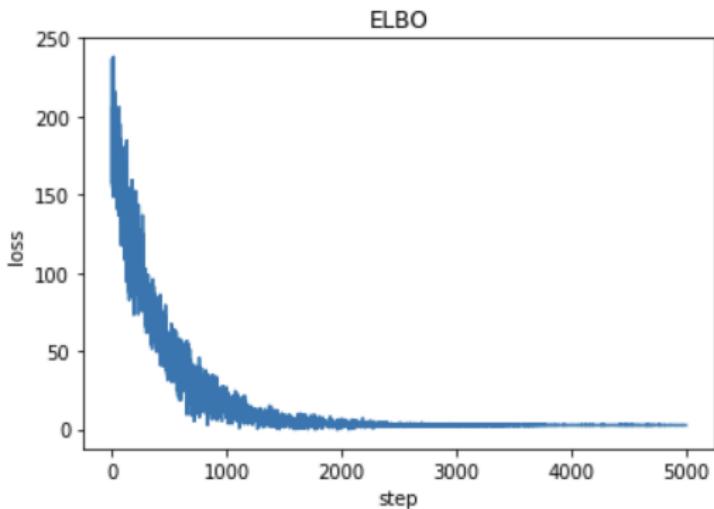
#Print results
print("P(Temperature|Sensor=18.0) = ")
print(dist.Normal(pyro.param("mean").item(), pyro.param("scale").item()))
print("")
```

Pyro's inference (http://pyro.ai/examples/intro_part_ii.html) :

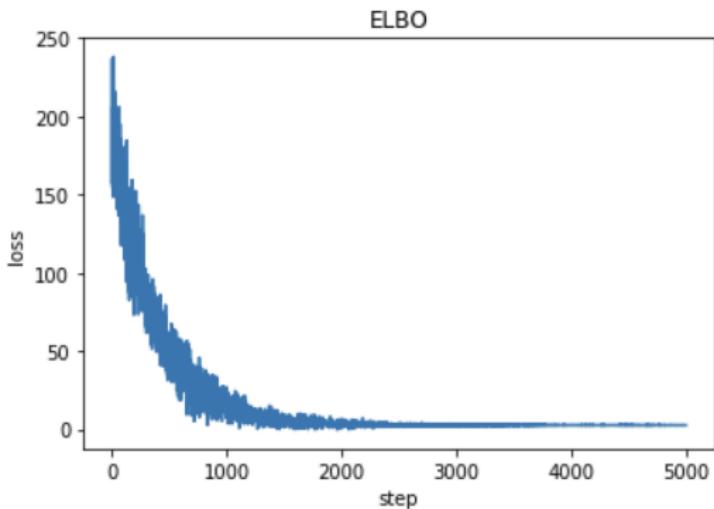
- We can introduce observations (e.g. sensor = 18.0).
- We can query the **posterior probability distribution**:

$$p(\text{temp}|\text{sensor} = 18) = \frac{p(\text{sensor} = 18|\text{temp})p(\text{temp})}{\int p(\text{sensor} = 18|\text{temp})p(\text{temp})d\text{temp}}$$

- Guide is an auxiliary method needed for inference (more details in the coming sessions).



```
P(Temperature|Sensor=18.0) =  
Normal(loc: 17.39859390258789, scale: 0.9089401960372925)
```



```
P(Temperature|Sensor=18.0) =  
Normal(loc: 17.39859390258789, scale: 0.9089401960372925)
```

Details on the inference method will be given on the following sessions.

Exercise 1: Change in the precision of the temperature sensor

- The precision of our temperature sensor is reflected in the **scale** of the Normal distribution of the **sensor** variable.
- What happens if we get a **more precise temperature sensor**? Assume it has a scale equal to 0.5.

Session1/students_PPLs_Intro.ipynb

```
In [ ]: #The observations
obs = {'sensor': torch.tensor([18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1])}

def model(obs):
    for i in range(0,obs['sensor'].shape[0]):
        temp = pyro.sample(f'temp_{i}', dist.Normal(15.0, 2.0))
        sensor = pyro.sample(f'sensor_{i}', dist.Normal(temp, 1.0), obs=obs['sensor'][i])
```

Pyro's inference (http://pyro.ai/examples/intro_part_ii.html) :

- What if we have a **bunch of observations**, $s = \{s_1, \dots, s_n\}$

```
In [ ]: #The observations
obs = {'sensor': torch.tensor([18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1])}

def model(obs):
    for i in range(0,obs['sensor'].shape[0]):
        temp = pyro.sample(f'temp_{i}', dist.Normal(15.0, 2.0))
        sensor = pyro.sample(f'sensor_{i}', dist.Normal(temp, 1.0), obs=obs['sensor'][i])
```

Pyro's inference (http://pyro.ai/examples/intro_part_ii.html) :

- What if we have a **bunch of observations**, $s = \{s_1, \dots, s_n\}$
- A random variable is created for each observation (using a **for-loop**).

```
In [ ]: #The observations
obs = {'sensor': torch.tensor([18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1])}

def model(obs):
    for i in range(0,obs['sensor'].shape[0]):
        temp = pyro.sample(f'temp_{i}', dist.Normal(15.0, 2.0))
        sensor = pyro.sample(f'sensor_{i}', dist.Normal(temp, 1.0), obs=obs['sensor'][i])
```

Pyro's inference (http://pyro.ai/examples/intro_part_ii.html) :

- What if we do not know the average temperature?

In [28]:

```
#The observations
obs = {'sensor': torch.tensor([18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1])}

def model(obs):
    mean_temp = pyro.param('mean_temp', torch.tensor(15.0))
    for i in range(0,obs['sensor'].shape[0]):
        temp = pyro.sample(f'temp_{i}', dist.Normal(mean_temp, 2.0))
        sensor = pyro.sample(f'sensor_{i}', dist.Normal(temp, 1.0), obs=obs['sensor'][i])
```

Pyro's inference (http://pyro.ai/examples/intro_part_ii.html) :

- What if we do not know the average temperature?
- We can introduce a **parameter** using **pyro.param** construct.

```
In [32]: #Run inference
    svi(model, guide, obs, num_steps=1000)

#Print results
print("Estimated Mean Temperature")
print(pyro.param("mean_temp").item())
```

```
Estimated Mean Temperature
19.129146575927734
```

Pyro's inference (http://pyro.ai/examples/intro_part_ii.html) :

- And **learn the parameter** with the same general inference algorithm.

$$\mu_t = \arg \max_{\mu} \ln p(s_1, \dots, s_n | \mu)$$

- Details about the inference algorithm will be given in the next sessions.

```
In [28]: #The observations
obs = {'sensor': torch.tensor([18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1])}

def model(obs):
    mean_temp = pyro.param('mean_temp', torch.tensor(15.0))
    for i in range(0,obs['sensor'].shape[0]):
        temp = pyro.sample(f'temp_{i}', dist.Normal(mean_temp, 2.0))
        sensor = pyro.sample(f'sensor_{i}', dist.Normal(temp, 1.0), obs=obs['sensor'][i])
```

Pyro's inference (http://pyro.ai/examples/intro_part_ii.html) :

- What if we want to capture uncertainty about the estimation of the average temperature?

```
In [176]: #The observations
obs = {'sensor': torch.tensor([18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1])}

def model(obs):
    mean_temp = pyro.sample('mean_temp', dist.Normal(15.0, 2.0))
    for i in range(obs['sensor'].shape[0]):
        temp = pyro.sample(f'temp_{i}', dist.Normal(mean_temp, 2.0))
        sensor = pyro.sample(f'sensor_{i}', dist.Normal(temp, 1.0), obs=obs['sensor'][i])
```

Pyro's inference (http://pyro.ai/examples/intro_part_ii.html) :

- What if we want to capture uncertainty about the estimation of the average temperature?
- We can model this parameter with a random variable.

```
In [162]: import time

#Run inference
start = time.time()
svi(model, guide, obs, num_steps=1000)

#print results
print("P(mean_temp|Sensor=[18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1]) =")
print(dist.Normal(pyro.param("mean").item(), pyro.param("scale").item()))
print("")
end = time.time()
print(f"\{(end - start)\} seconds")

P(mean_temp|Sensor=[18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1]) =
Normal(loc: 19.199871063232422, scale: 0.6046891212463379)

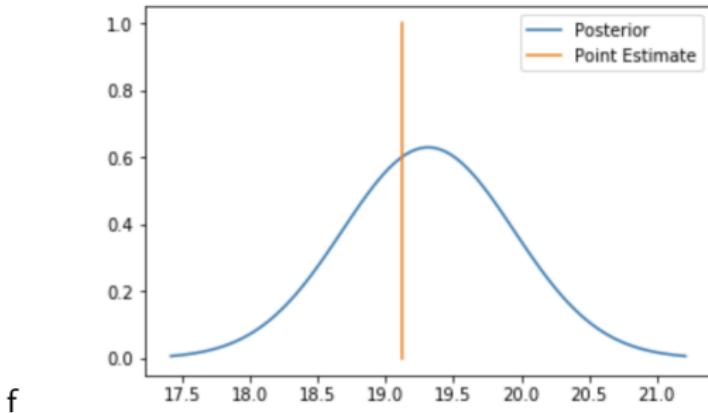
10.298431873321533 seconds
```

Pyro's inference (http://pyro.ai/examples/intro_part_ii.html) :

- And **learn the distribution** with the same general inference algorithm.

$$p(\mu_t | s_1, \dots, s_{10})$$

- Details about the inference algorithm will be given in the next sessions.



Pyro's inference (http://pyro.ai/examples/intro_part_ii.html) :

- And **learn the distribution** with the same general inference algorithm.

$$p(\mu_t | s_1, \dots, s_{10})$$

- Details about the inference algorithm will be given in the next sessions.

Defining Conditional Independences in Pyro

```
In [176]: #The observations
obs = {'sensor': torch.tensor([18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1])}

def model(obs):
    mean_temp = pyro.sample('mean_temp', dist.Normal(15.0, 2.0))
    for i in range(obs['sensor'].shape[0]):
        temp = pyro.sample(f'temp_{i}', dist.Normal(mean_temp, 2.0))
        sensor = pyro.sample(f'sensor_{i}', dist.Normal(temp, 1.0), obs=obs['sensor'][i])
```

Pyro's cond. independences (http://pyro.ai/examples/svi_part_ii.html) :

- Sensor variables are **independent** given temperature mean, μ_t .

$$p(s_1, t_1, \dots, s_{10}, t_{10} | \mu_t) = \prod_{i=1}^{10} p(s_i, t_i | \mu_t)$$

```
In [37]: #The observations
obs = {'sensor': torch.tensor([18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1])}

def model(obs):
    mean_temp = pyro.sample('mean_temp', dist.Normal(15.0, 2.0))
    with pyro.plate('a', obs['sensor'].shape[0]):
        temp = pyro.sample('temp', dist.Normal(mean_temp, 2.0))
        sensor = pyro.sample('sensor', dist.Normal(temp, 1.0), obs=obs['sensor'])
```

Pyro's cond. independences (http://pyro.ai/examples/svi_part_ii.html) :

- Sensor variables are **independent** given temperature mean, μ_t .

$$p(s_1, t_1, \dots, s_{10}, t_{10} | \mu_t) = \prod_{i=1}^{10} p(s_i, t_i | \mu_t)$$

- We can use **Pyro's plate construct** to introduce this independence.

```
In [165]: #Run inference
start = time.time()
svi(model, guide, obs, num_steps=1000)

#Print results
print("P(mean_temp|Sensor=[18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1]) =")
print(dist.Normal(pyro.param("mean").item(), pyro.param("scale").item()))
print("")
end = time.time()
print(f"{(end - start)} seconds")

P(mean_temp|Sensor=[18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1]) =
Normal(loc: 19.300748825073242, scale: 0.6379732489585876)

2.81210994720459 seconds
```

Pyro's cond. independences (http://pyro.ai/examples/svi_part_ii.html) :

- We get large gains in efficiency due to **vectorized operations**.
- Execution time without *plate* is over 10s.

- **Exercise 2:** The role of the **number of observations** in learning.
- **Exercise 3:** The role of the **prior distribution** in learning.

`Day1/students_PPLs_Intro.ipynb`

Toy Example: Ice-cream shop



Defining Machine Learning models with PPLs:

- We have an ice-cream shop and we record the ice-cream sales and the average temperature of the day.
- We know temperature affects the sales of ice-creams.
- We want to precisely find out how temperature affects ice-cream sales.

```
In [49]: #The observations
obs = {'sensor': torch.tensor([18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1]),
       'sales': torch.tensor([46., 47., 49., 44., 50., 54., 51., 52., 49., 53.])}

def model(obs):
    mean_temp = pyro.sample('mean_temp', dist.Normal(15.0, 2.0))

    with pyro.plate('a', obs['sensor'].shape[0]):
        temp = pyro.sample('temp', dist.Normal(mean_temp, 2.0))
        sensor = pyro.sample('sensor', dist.Normal(temp, 1.0), obs=obs['sensor'])
```

Ice-cream Shop Model:

- We have observations from temperature and sales.

```
In [43]: #The observations
obs = {'sensor': torch.tensor([18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5, 20.1]),
       'sales': torch.tensor([46., 47., 49., 44., 50., 54., 51., 52., 49., 53.])}

def model(obs):
    mean_temp = pyro.sample('mean_temp', dist.Normal(15.0, 2.0))

    with pyro.plate('a', obs['sensor'].shape[0]):
        temp = pyro.sample('temp', dist.Normal(mean_temp, 2.0))
        sensor = pyro.sample('sensor', dist.Normal(temp, 1.0), obs=obs['sensor'])
        sales = pyro.sample('sales', dist.Poisson(??????), obs=obs['sales'])
```

Ice-cream Shop Model:

- We have observations from temperature and sales.
- Sales are modeled with a **Poisson distribution**.

```
In [201]: #The observations
obs = {'sensor': torch.tensor([18., 18.7, 19.2, 17.8, 20.3, 22.4, 20.3, 21.2, 19.5]),
       'sales': torch.tensor([46., 47., 49., 44., 50., 54., 51., 52., 49., 53.])}

def model(obs):
    mean_temp = pyro.sample('mean_temp', dist.Normal(15.0, 2.0))
    alpha = pyro.sample('alpha', dist.Normal(0.0, 100.0))
    beta = pyro.sample('beta', dist.Normal(0.0, 100.0))

    with pyro.plate('a', obs['sensor'].shape[0]):
        temp = pyro.sample('temp', dist.Normal(mean_temp, 2.0))
        sensor = pyro.sample('sensor', dist.Normal(temp, 1.0), obs=obs['sensor'])
        rate = torch.max(torch.tensor(0.001), alpha + beta*temp)
        sales = pyro.sample('sales', dist.Poisson(rate), obs=obs['sales'])
```

Ice-cream Shop Model:

- We have observations from temperature and sales.
 - Sales are modeled with a **Poisson distribution**.
 - The rate of the Poisson **linearly depends of the real temperature**.

```
In [48]: #Run inference
svi(model, guide, obs, num_steps=1000)

#Print results
print("Posterior Temperature Mean")
print(dist.Normal(pyro.param("mean").item(), pyro.param("scale").item()))
print("")
print("Posterior Alpha")
print(dist.Normal(pyro.param("alpha_mean").item(), pyro.param("alpha_scale").item()))
print("")
print("Posterior Beta")
print(dist.Normal(pyro.param("beta_mean").item(), pyro.param("beta_scale").item()))

Posterior Temperature Mean
Normal(loc: 19.311052322387695, scale: 0.6258021593093872)

Posterior Alpha
Normal(loc: 19.773971557617188, scale: 1.8541947603225708)

Posterior Beta
Normal(loc: 1.5178951025009155, scale: 0.1155082955956459)
```

Ice-cream Shop Model:

- We run the **(variational) inference engine** and get the results.

```
In [48]: #Run inference
svi(model, guide, obs, num_steps=1000)

#Print results
print("Posterior Temperature Mean")
print(dist.Normal(pyro.param("mean").item(), pyro.param("scale").item()))
print("")
print("Posterior Alpha")
print(dist.Normal(pyro.param("alpha_mean").item(), pyro.param("alpha_scale").item()))
print("")
print("Posterior Beta")
print(dist.Normal(pyro.param("beta_mean").item(), pyro.param("beta_scale").item()))

Posterior Temperature Mean
Normal(loc: 19.311052322387695, scale: 0.6258021593093872)

Posterior Alpha
Normal(loc: 19.773971557617188, scale: 1.8541947603225708)

Posterior Beta
Normal(loc: 1.5178951025009155, scale: 0.1155082955956459)
```

Ice-cream Shop Model:

- We run the **(variational) inference engine** and get the results.
- With PPLs, we only care about modeling, **not about the low-level details** of the machine-learning solver.

Exercise 4: Introduce Humidity in the Icecream shop model

- Assume we also have a bunch of humidity sensor measurements.
- Assume the sales are also linearly influenced by the humidity.
- **Extend the above model in order to integrate all of that.**

`Day1/students_PPLs_Intro.ipynb`